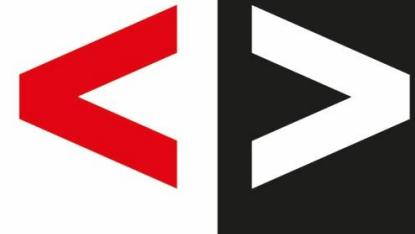


Módulo 6 MDA - Advanced Analytics

Machine Learning

Sesión 1 - Clasificación

Sesión 2 - Clusterización



José Juan Martínez Carrascosa

Jose J. Martinez

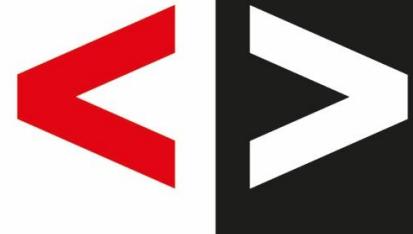
jjmcarrascosa@gmail.com

Índice (sesión teoría)

1. Sobre las dinámicas y tiempos de clase
2. **Inteligencia Artificial y Machine Learning**
 - a. Activación de conocimientos previos
 - b. Inteligencia Artificial
 - c. Negocio e IA (McKinsey 2017)
 - d. Data Science y Data Analytics. Estadística y ML.
 - e. Machine Learning. Deep Learning
3. **ML supervisado vs no supervisado**
4. ML supervisado: **Clasificación**
 - a. Por tipo de problema: binaria, multiclase, multietiqueta
 - b. Por tipo de algoritmo: regresión, modelos bayesianos, descenso por gradiente, SVM, árboles de decisión y boosting
 - c. Métricas
5. ML no supervisado: **Clusterización**
 - a. Por tipo: de partición, de densidad, de rejilla.
 - b. Métricas
6. Práctica 1: Clasificación binaria
7. Práctica 2: Clusterización



Sobre las dinámicas y tiempos de clase



Dinámicas de clase

FRONTAL



*Introducción a conceptos
nuevos, explicaciones
magistrales*

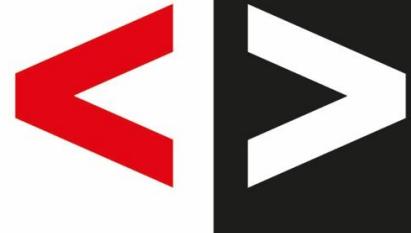
Teoría: 2 horas
Práctica: 45 minutos

GRUPOS



*Anticipación, deducción,
brainstorming*

Teoría: 1 hora
Práctica: 2h 15 minutos

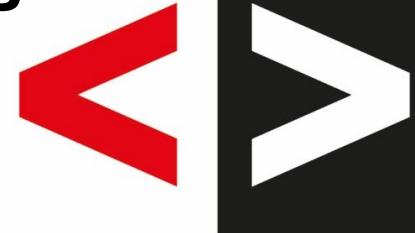


<ISDI>

Inteligencia Artificial y Machine Learning

Activación de conocimientos previos

<https://forms.gle/bn2KqyQRjaRTQCHF9>



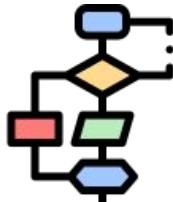
Inteligencia Artificial

“

La **Inteligencia Artificial** (IA) es la combinación de algoritmos alimentados con datos que cuyo propósito es el crear modelos estadísticos que presenten las mismas capacidades de ejecución que el ser humano

”

Algoritmos



+

Datos



=



Modelos estadísticos



Inteligencia Artificial

```

repeat until convergence {
     $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})$ 
     $\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$ 
}
 $J = \frac{1}{n} \sum_{i=1}^n (pred_i - y_i)^2$ 
 $J = \frac{1}{n} \sum_{i=1}^n (a_0 + a_1 \cdot x_i - y_i)^2$ 

```

Now,

$$\frac{\partial}{\partial \theta} J_\theta = \frac{\partial}{\partial \theta} \frac{1}{2n}$$



Therefore,

$$dZ^{[L]} = A^{[L]} - Y$$

$$dW^{[L]} = \frac{1}{m} dZ^{[L]} A^{[L]T}$$

$$db^{[L]} = \frac{1}{m} np.sum(dZ^{[L]}, axis = 1, keepdims = True)$$

$$dZ^{[L-1]} = W^{[L]T} dZ^{[L]} g^{[L]}(Z^{[L-1]})$$

$$dZ^{[1]} = W^{[1]T} dZ^{[2]} g^{[1]}(Z^{[1]})$$

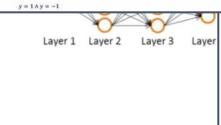
$$dW^{[1]} = \frac{1}{m} dZ^{[1]} A^{[1]T}$$

$$db^{[1]} = \frac{1}{m} np.sum(dZ^{[1]}, axis = 1, k$$

IA = Formulas + Datos

Given one training example (

Forward propagation:

$$\begin{aligned} a^{(1)} &= x \\ z^{(2)} &= \Theta^{(1)} a^{(1)} \\ \rightarrow a^{(2)} &= g(z^{(2)}) \quad (\text{add } a_0) \\ \rightarrow z^{(3)} &= \Theta^{(2)} a^{(2)} \\ \rightarrow a^{(3)} &= g(z^{(3)}) \quad (\text{add } a_0) \\ \rightarrow z^{(4)} &= \Theta^{(3)} a^{(3)} \\ \rightarrow a^{(4)} &= h_\theta(x) = g(z^{(4)}) \end{aligned}$$


Given one training example (

Policy Iteration
We can see that the optimal policy π^* is not necessarily the same as the current policy π . This is because the value function V^π is not necessarily the same as the true value function V^* . We can use the Bellman equation to update the policy π based on the new value function V^π .

K-Nearest Neighbor
 $D(x_i, x_j) = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$

Support Vector Machines
 $f(x) = \max_{i=1}^n w_i x_i + b$

Backpropagation
 $d\sigma_i(x) = \sigma_i(x) \cdot (1 - \sigma_i(x))$

Logistic Regression
 $\text{Odds Ratio} = \log \frac{P(y)}{1 - P(y)}$

Principal Components Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

Logistic Regression

Gradient Descent

Principle Component Analysis

Neural Networks

Support Vector Machines

Backpropagation

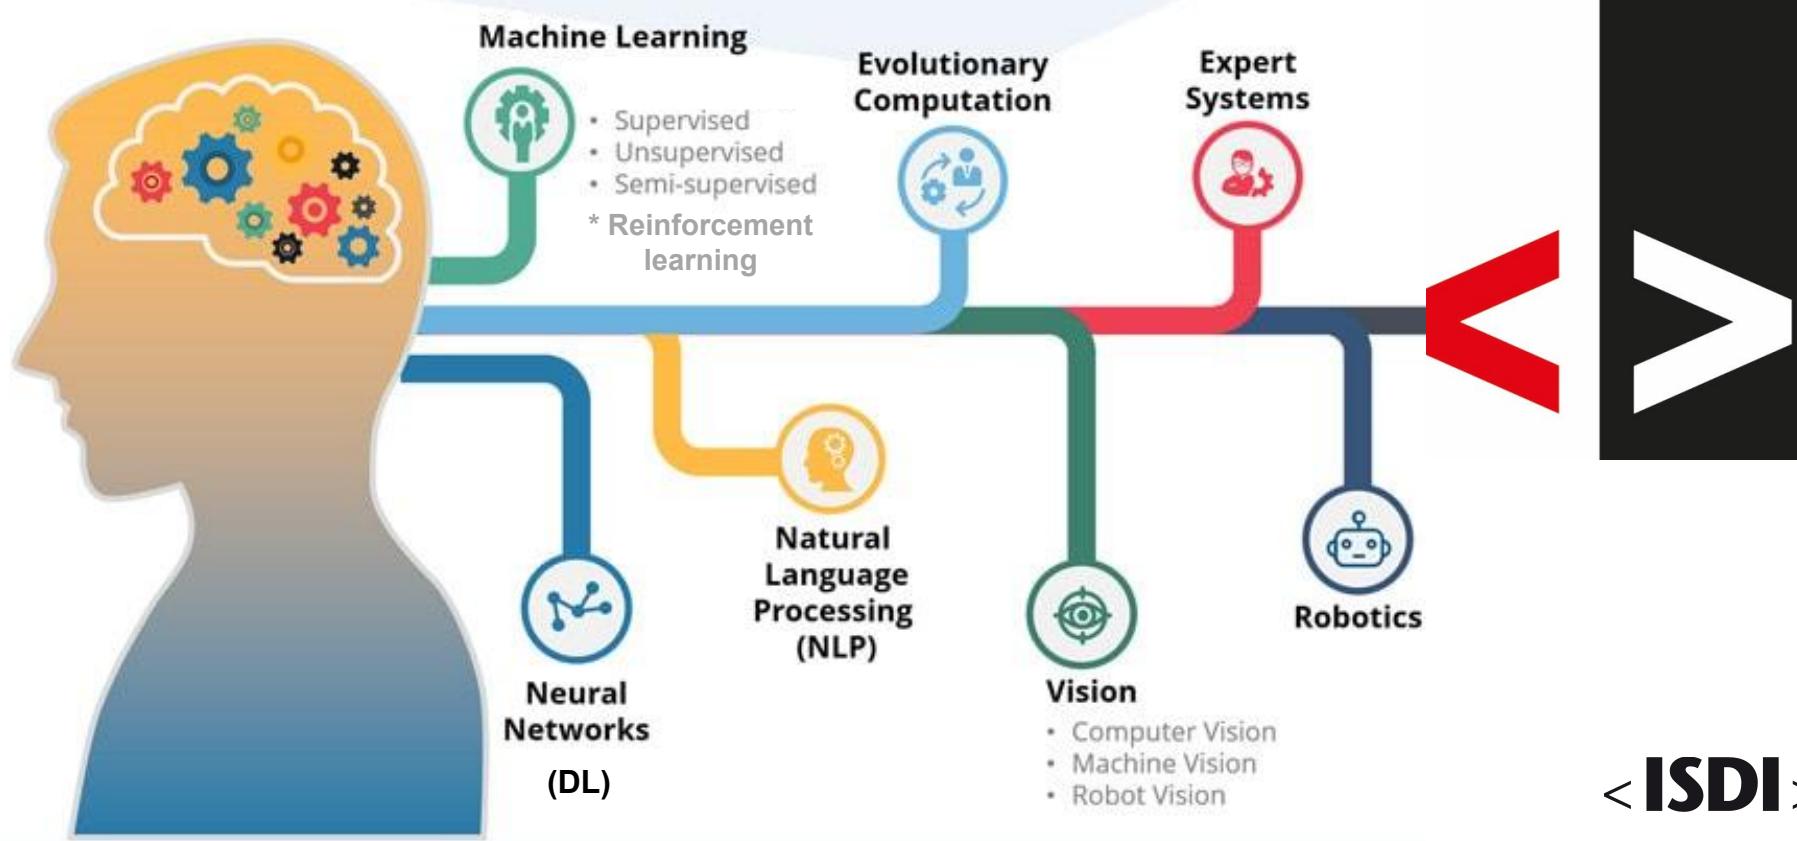
Logistic Regression

Gradient Descent

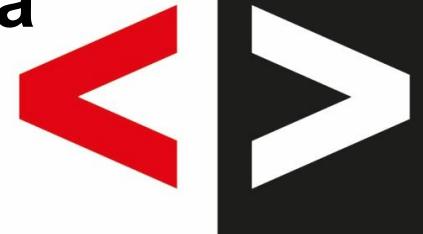
Principle Component Analysis

Neural Networks

Inteligencia Artificial



Potencial de automatización por area de negocio IA



Potencial de automatización por áreas

Time spent in US occupations, %



Technical feasibility, % of time spent on activities that can be automated by adapting currently demonstrated technology

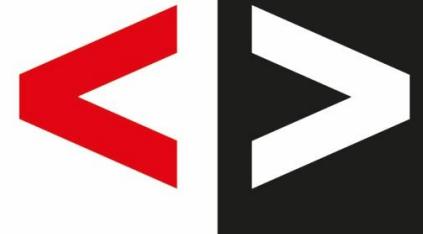


Wholesale trade
Finance and insurance
Arts, entertainment, and recreation
Real estate
Administrative
Healthcare and social assistance
Information
Professional
Management
Educational services



McKinsey, 2017

Data Science y Data Analytics



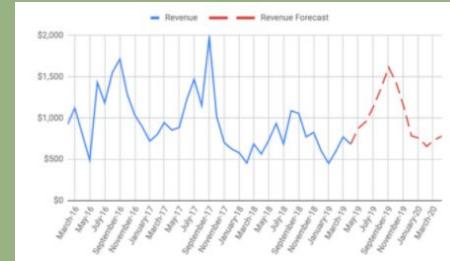
Analítica
descriptiva

Business
Intelligence

Analítica predictiva

Data
Science

Analítica prescriptiva



PASADO

PRESENTE

FUTURO

Analítica descriptiva

Business Intelligence

Paneles de control y reportes

KPI

Clusterización de clientes

Monitorización de consumos

Cruce de datos, agregaciones, proyecciones



Analítica predictiva

Data Science

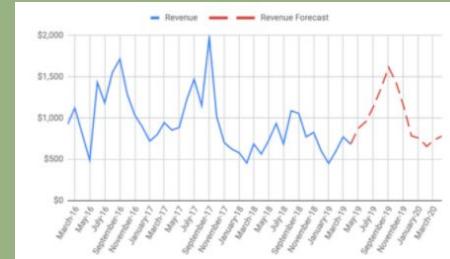
Forecasting de consumos, ventas, ambulatorización, contagios...

Analítica prescriptiva

Retención de clientes – CHURN

Automatización de toma de decisiones

Alertas sobre tendencias futuras

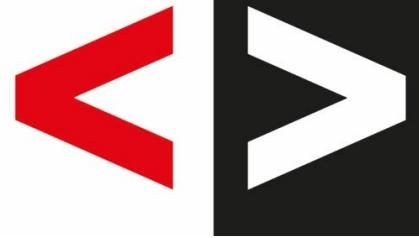


PASADO

PRESENTE

FUTURO

Estadística y ML



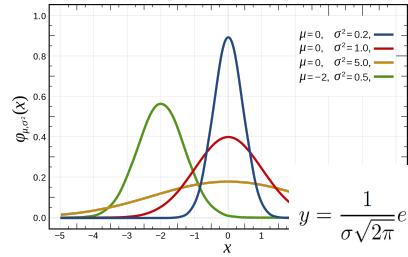
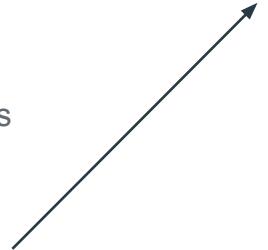
Estamos rodeados de sucesos complejos y no vemos correlaciones estadísticas que ML si ve

Example of Normal (Laplace-Gauss) Distribution in the gym



ESTADISTICA

La **estadística** consiste en **conocer y aplicar** las mejores **formulas** que modelen los datos



$$y = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

μ = Mean

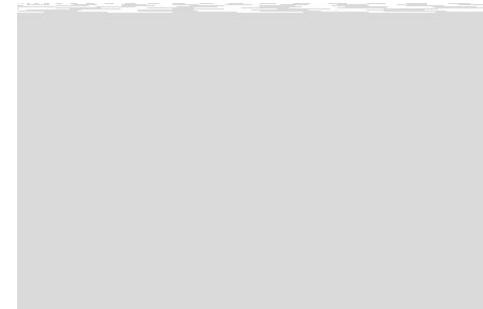
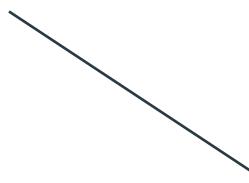
σ = Standard Deviation

$\pi \approx 3.14159 \dots$

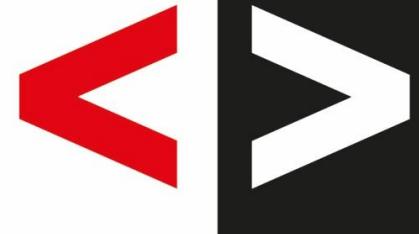
$e \approx 2.71828 \dots$

MACHINE LEARNING

ML se basa en los **datos** y automatiza ese modelado y aplicación de fórmulas

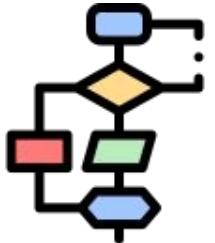


Machine Learning



ML, regresión (entrenamiento)

Algoritmo:
Regresión lineal



Datos anotados
(var. dep e
indep)

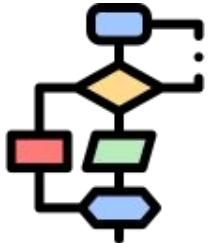


MES	DIA	OUTPUT
Junio	16	0.7
Diciembre	21	21
...

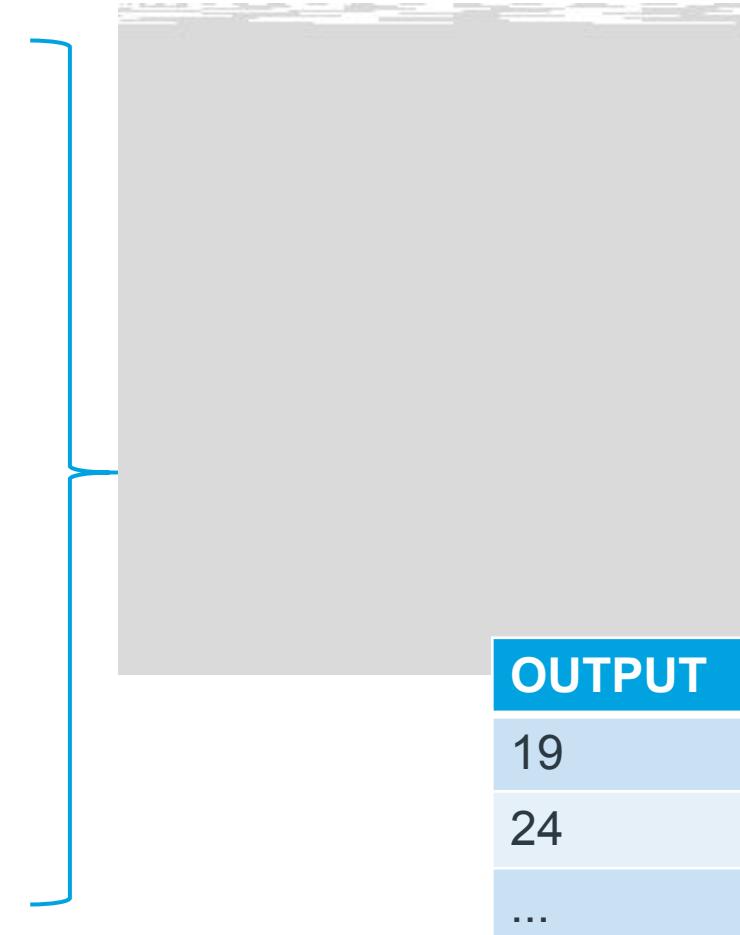
ML, regresión (inferencia)

Algoritmo:
Regresión lineal

Datos anotados
(var. dep e
indep)



MES	DIA
Enero	6
Marzo	11
...	...



Deep Learning

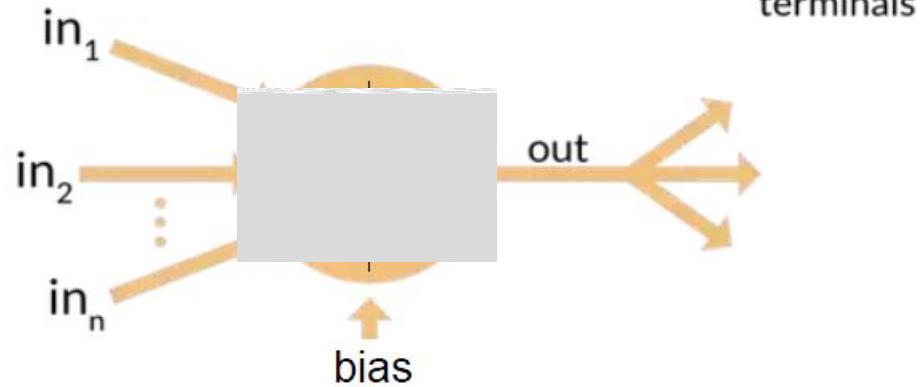
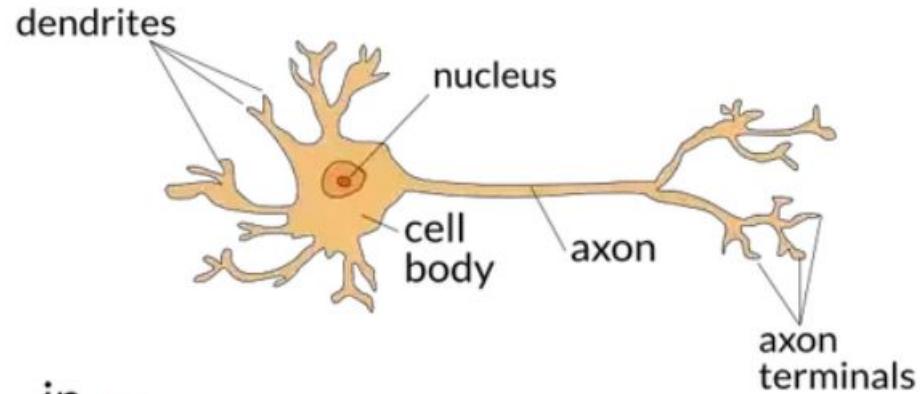
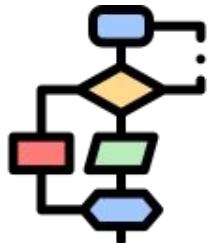


ML, regresión (entrenamiento)

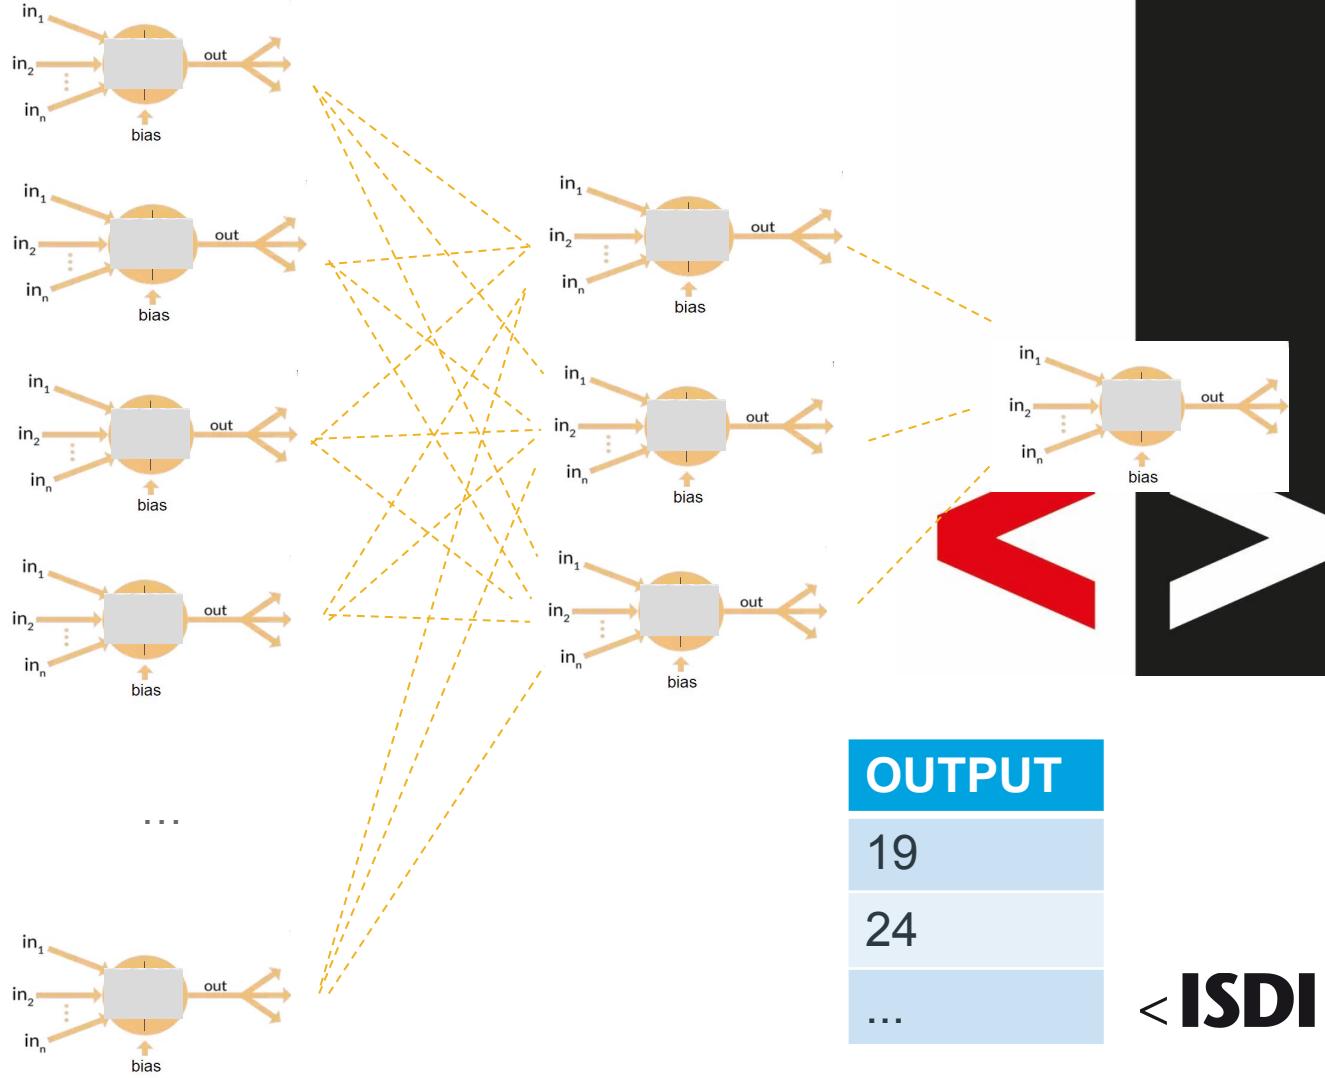
Algoritmo:
Regresión lineal

Datos anotados
(var. dep e
indep)

MES	DIA	OUTPUT
Junio	16	0.7
Diciembre	21	21
...



MES	DIA
Junio	16
Diciembre	21
...	...



OUTPUT
19
24
...

<ISDI>

Final del resumen





Artificial Intelligence

Any technique which enable computers to mimic intelligence

ARTIFICIAL INTELLIGENCE

MACHINE LEARNING

NLP
CV
....

DEEP LEARNING



Machine Learning

Subset of AI techniques which use statistical methods to enable machines to improve a result with experience

Deep Learning

Subset of machine learning that drives un/supervised learning from data that is unstructured or unlabeled



Natural Language Processing and Computer Vision

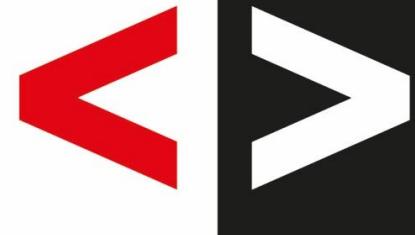
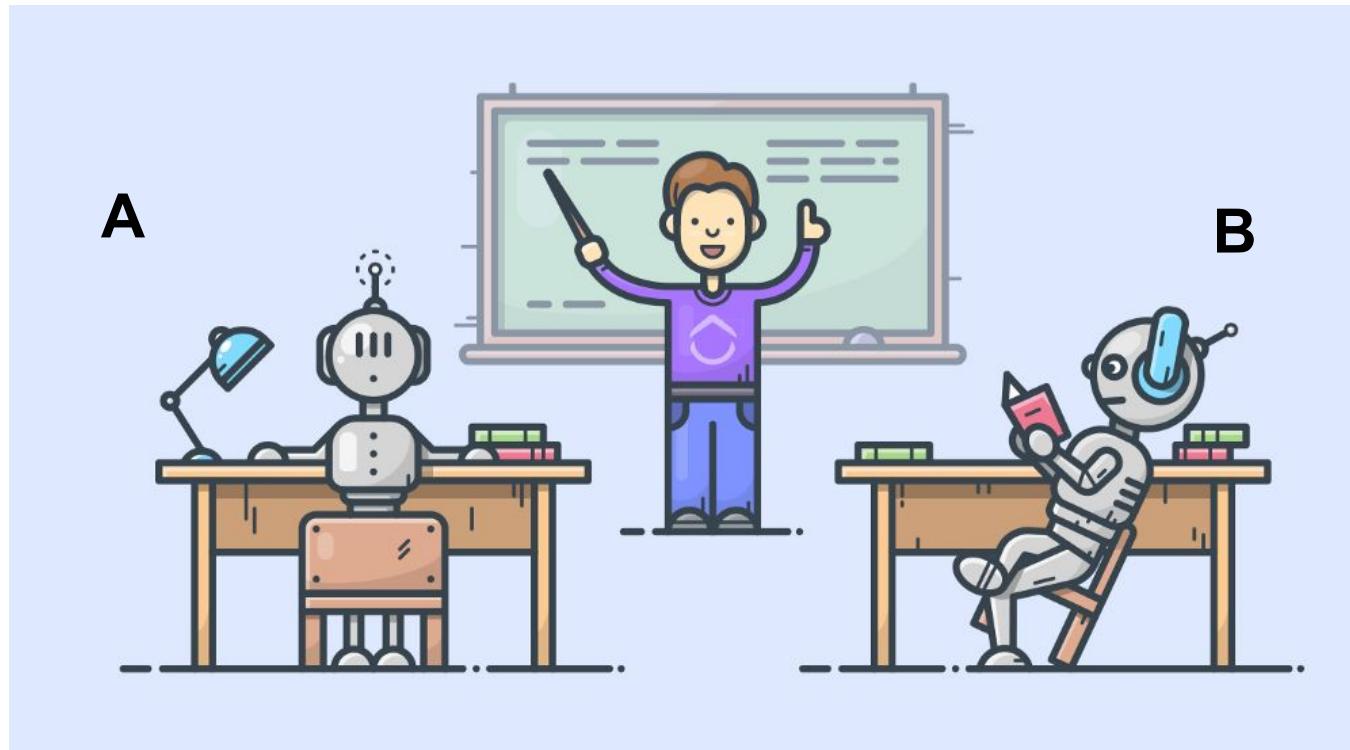
Programming that allows computers to understand human language or images



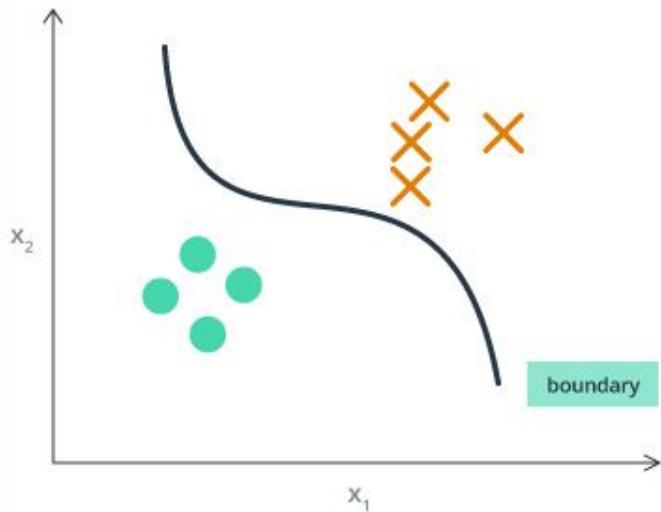
Machine Learning supervisado y no supervisado



¿Quien está aprendiendo de forma no supervisada?

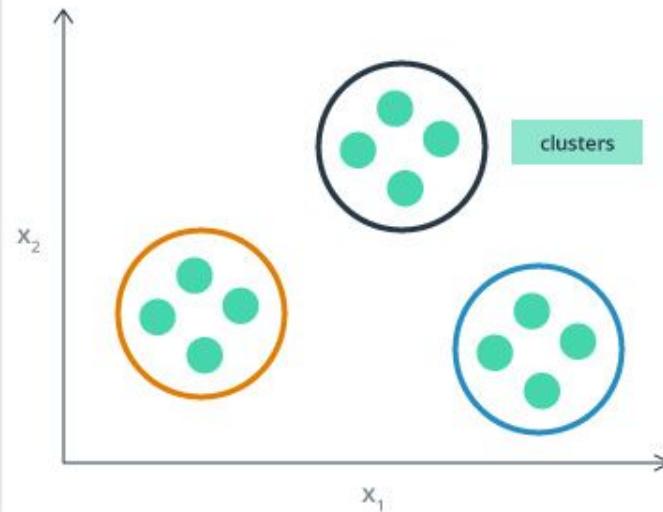


Supervised learning

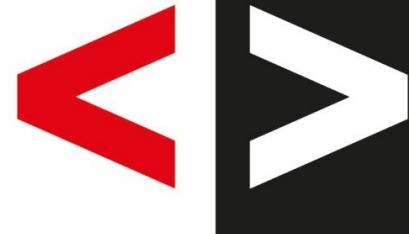


Ej: Regresión, clasificación, etc

Unsupervised learning

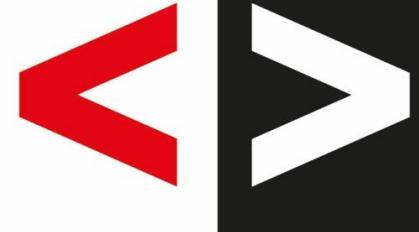


Ej: Clusterización, reducción de dimensionalidad, modelos de lengua en NLP, etc

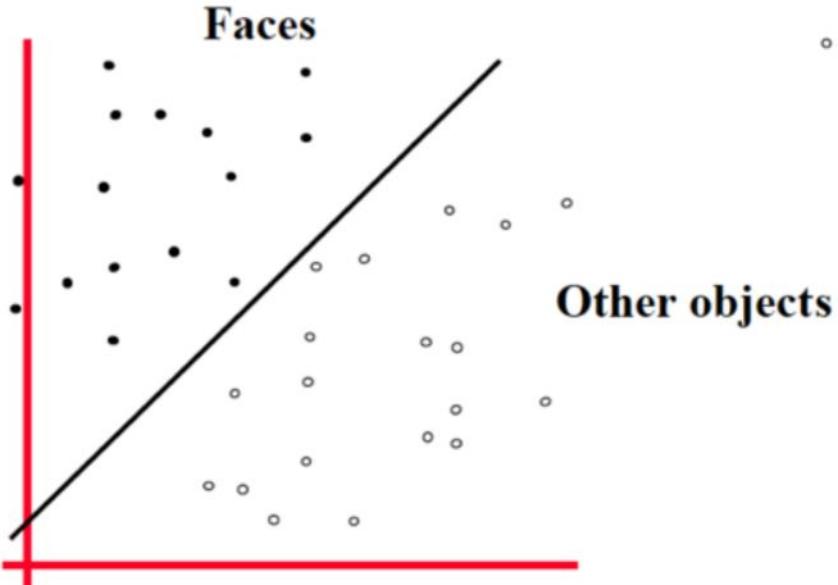


Aprendizaje supervisado	Aprendizaje no supervisado
Datos <i>etiquetados</i> a priori	Datos <i>no etiquetados</i>
Output <i>conocido</i>	Output <i>desconocido</i>
Enfoque <i>predictivo</i>	Enfoque <i>descriptivo</i>
<i>Regresión, clasificación</i>	<i>Clusterización, asociación, modelos de lengua</i>
<i>Más complejo de entrenar, necesita etiquetado humano</i>	<i>Menos complejo</i>
<i>Muchos casos de uso</i>	Menos casos de uso pero muy relevante ultimamente en NLP con modelos de lengua como Google BERT

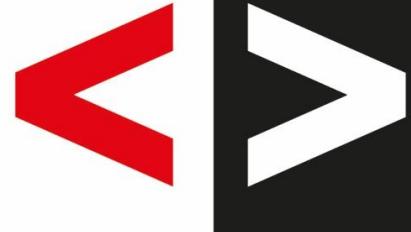
Machine Learning supervisado: Clasificación



Clasificación vs Regresión



- denotes +1 (faces)
- denotes -1 (other)



Idea: Find a separating hyperplane (line in this case)

Tipos de clasificación

Binary
Classification



- Spam
- Not spam

Multiclass
Classification

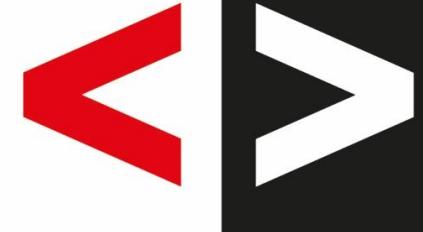


- Dog
- Cat
- Horse
- Fish
- Bird
- ...

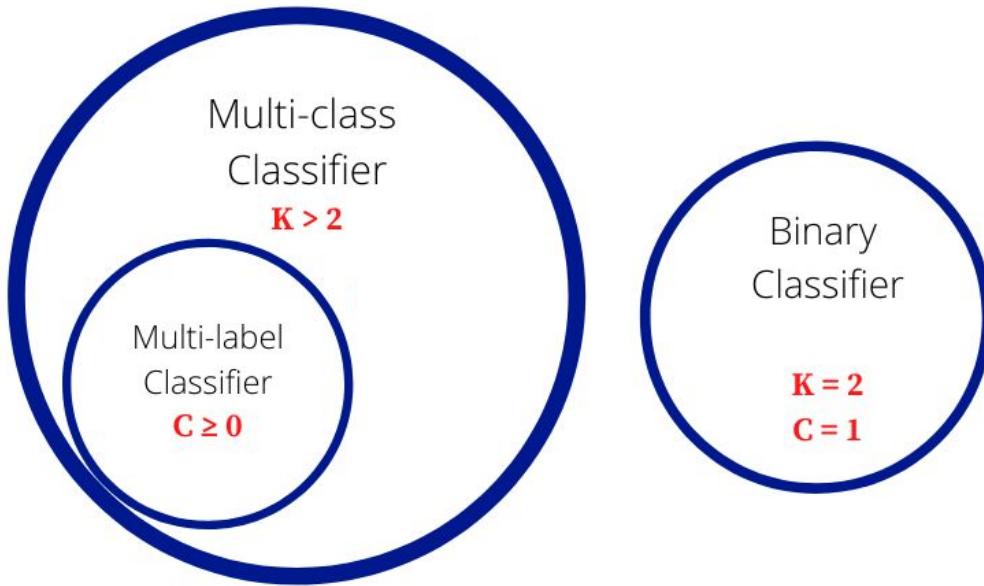
Multi-label
Classification



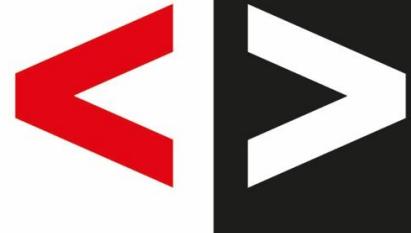
- Dog
- Cat
- Horse
- Fish
- Bird



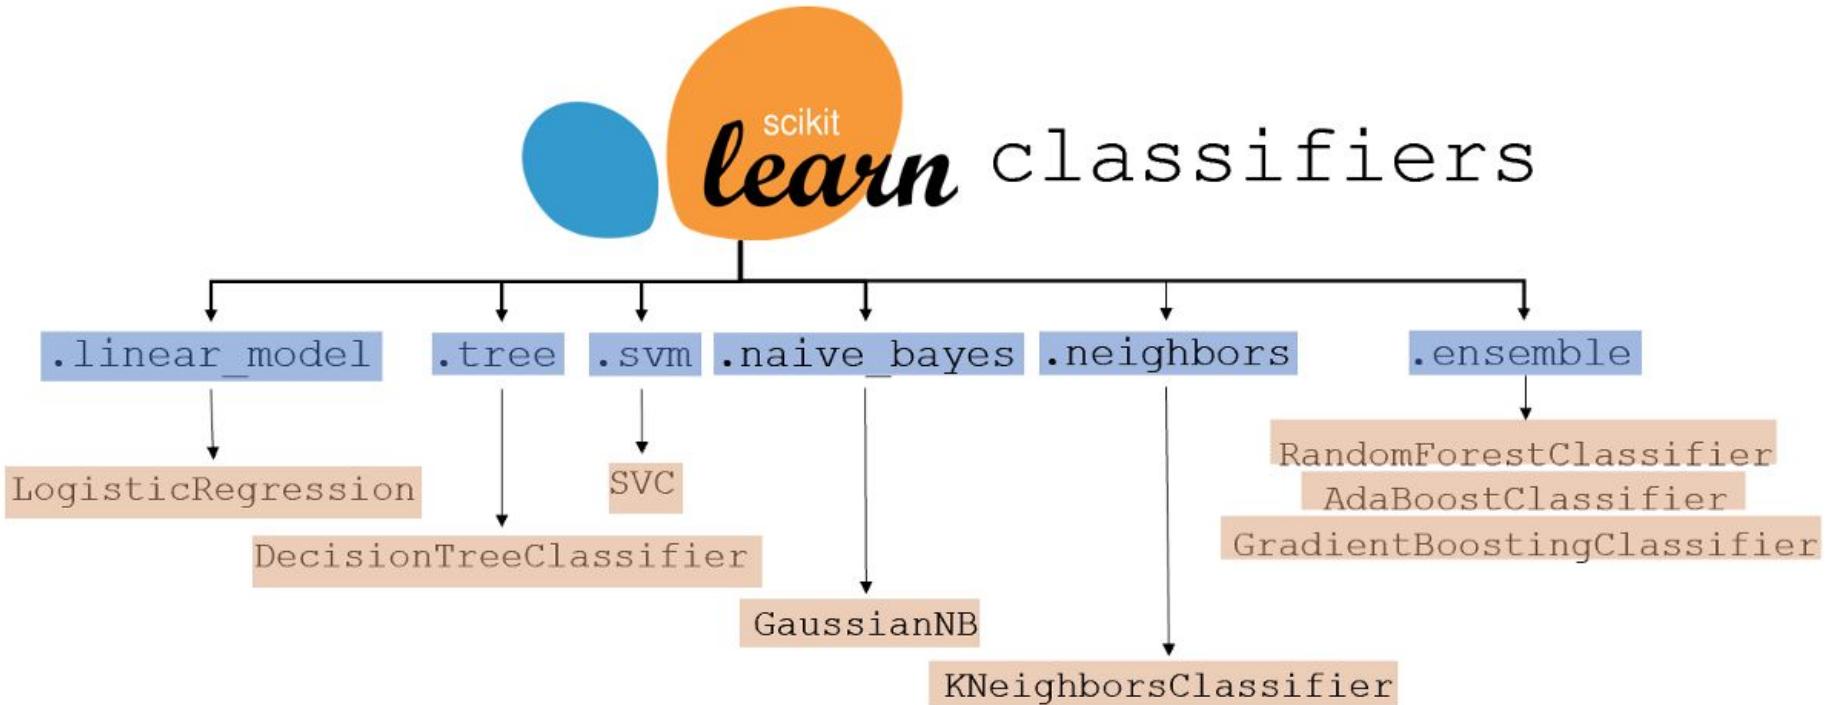
Tipos de clasificación



K = Total number of classes in the problem statement
 C = Number of classes an item maybe assigned to



Tipos de algoritmos

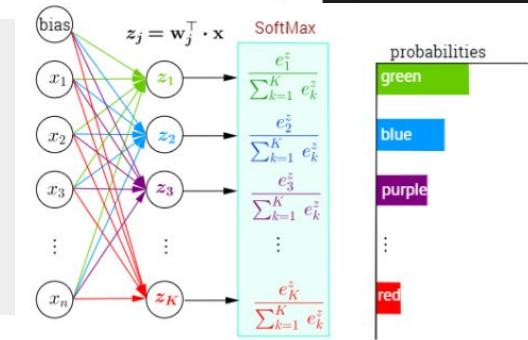
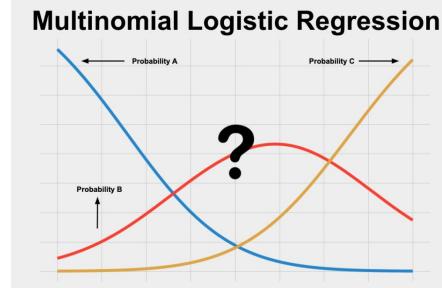
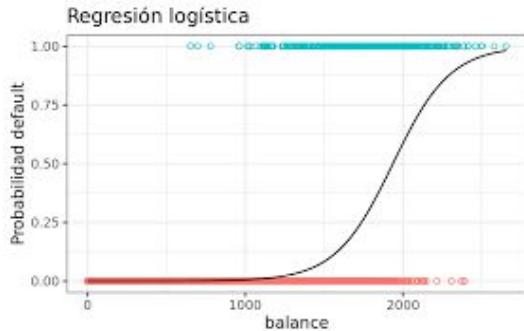


https://scikit-learn.org/stable/supervised_learning.html

Tipos de algoritmos

No todos los algoritmos soportan las 3 modalidades (binaria, multiclase, multietiqueta).

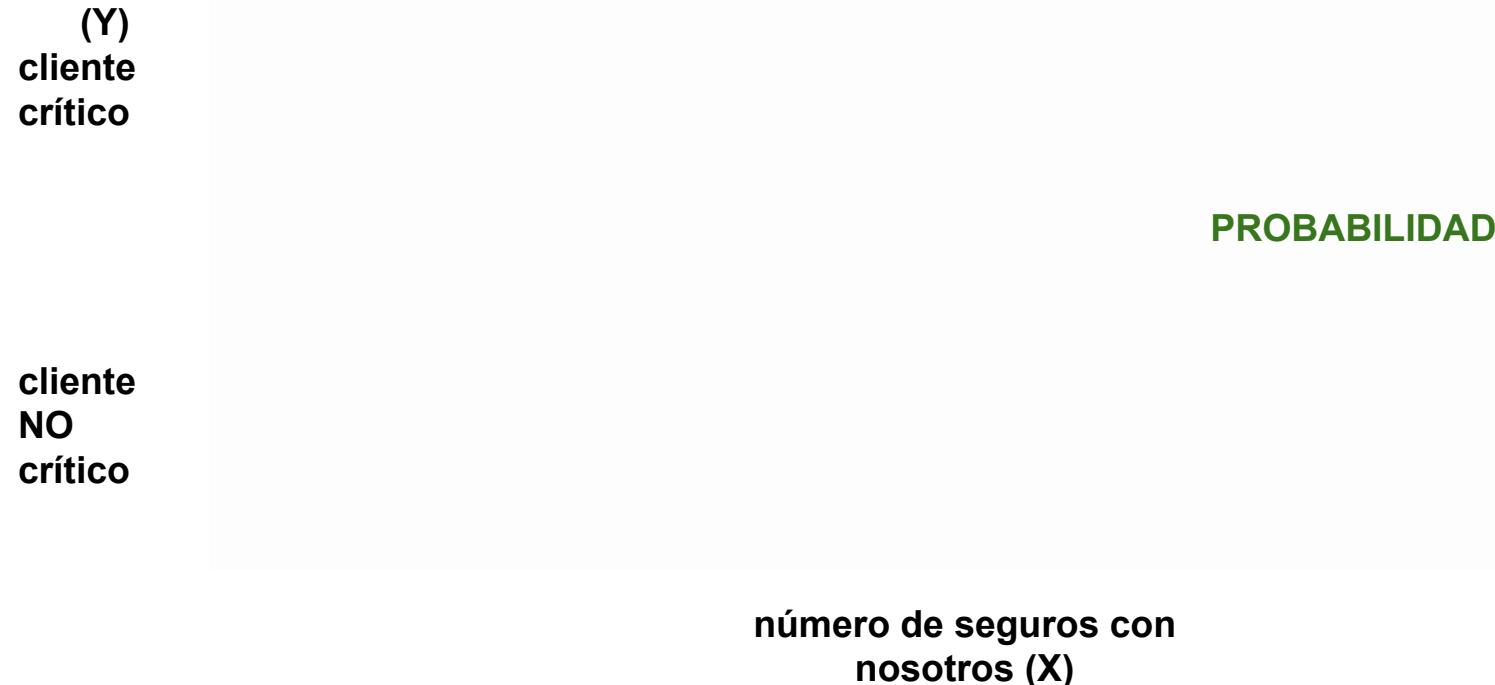
Aunque muchos de ellos tienen adaptaciones.



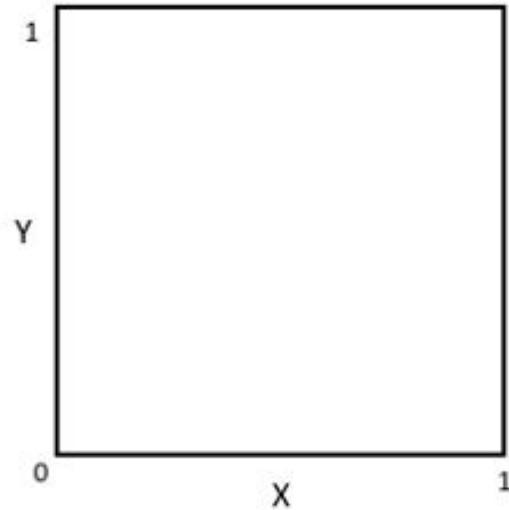
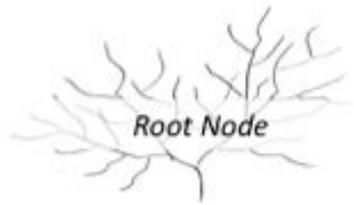
Para saber qué algoritmos están disponibles, acceder a:

<https://scikit-learn.org/stable/modules/multiclass.html>

Modelos lineares: Regresión logística



Árboles de decisión (Decision Trees)

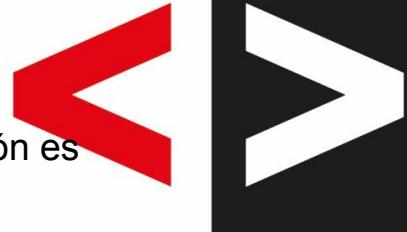


Ensemble

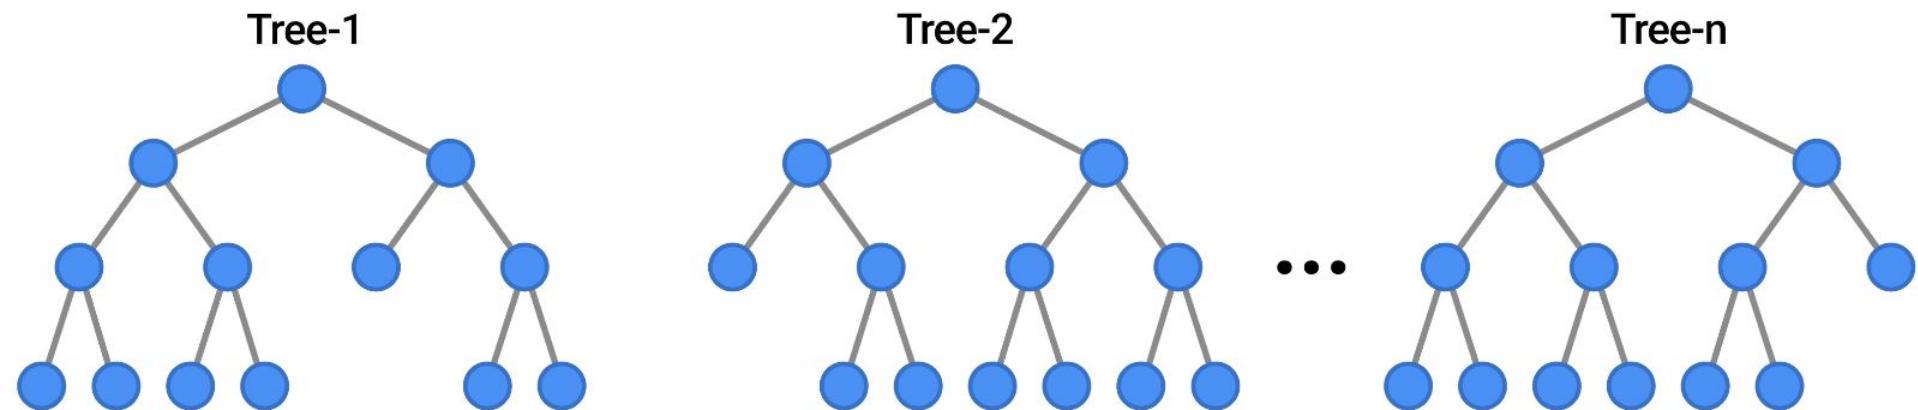
Algoritmos de aprendizaje para obtener un mejor rendimiento predictivo que el que se podría obtener de cualquiera de los algoritmos de aprendizaje constituyentes
Se distinguen dos familias de métodos de conjunto:

En los métodos de votación (bagging): el principio impulsor es construir varios estimadores de forma independiente y luego promediar sus predicciones. Ej: RandomForestTrees.

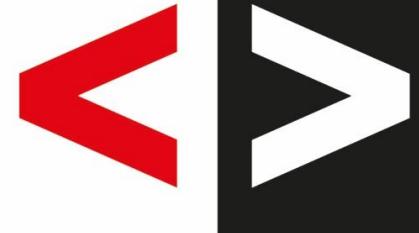
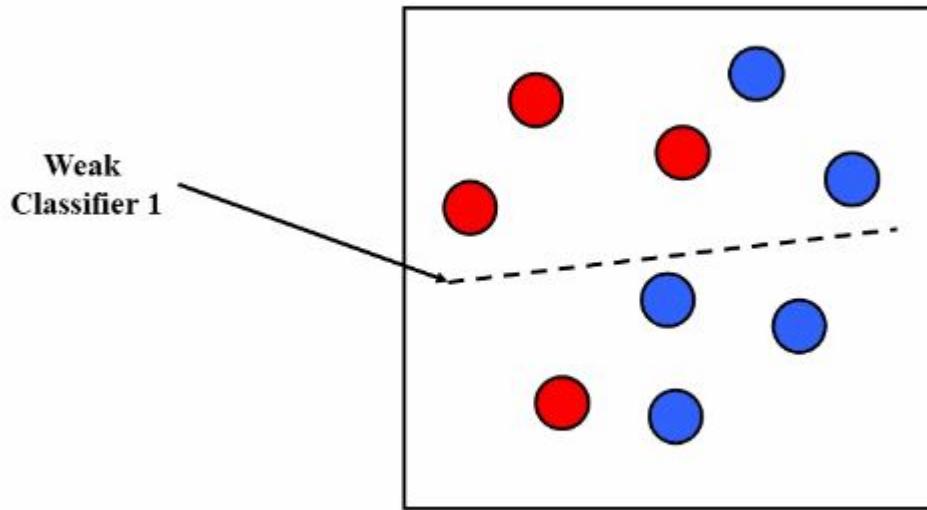
Los métodos de impulso (boosting), los estimadores base se construyen secuencialmente y se intenta reducir el sesgo del estimador acumulado. La motivación es combinar varios modelos débiles para producir un modelo mejor. Ejemplo: AdaBoost, XGBoost.



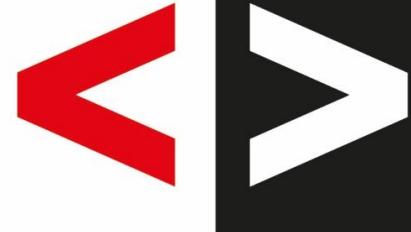
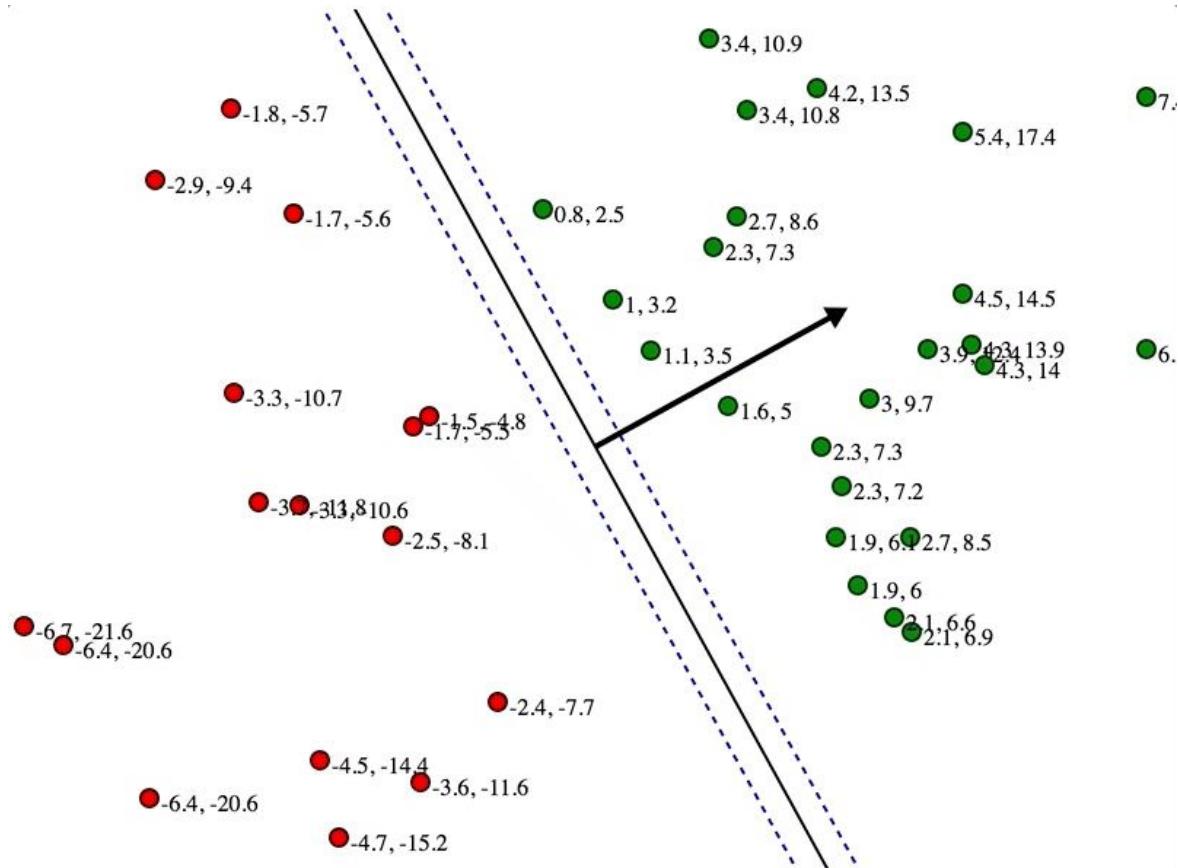
Ensemble - Bagging (votación). Random Forest



Ensemble - Boosting (impulso) - Adaboost



Support Vector Machines



Modelos Bayesianos

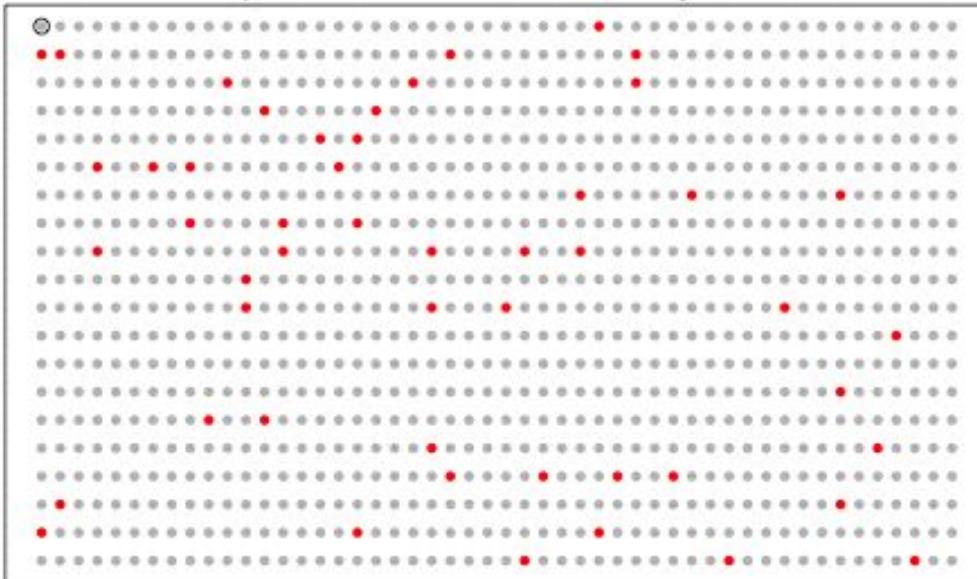
Modelos basados en el Teorema de Bayes, que calcula la probabilidad de darse una clase A teniendo una observación B, en base a la probabilidad de haberse visto B y ser una clase A.

CLASSIFIER

"Gaussian" because this is a normal distribution

This is our prior belief

$$P(\text{class} \mid \text{data}) = \frac{P(\text{data} \mid \text{class}) \times p(\text{class})}{P(\text{data})}$$



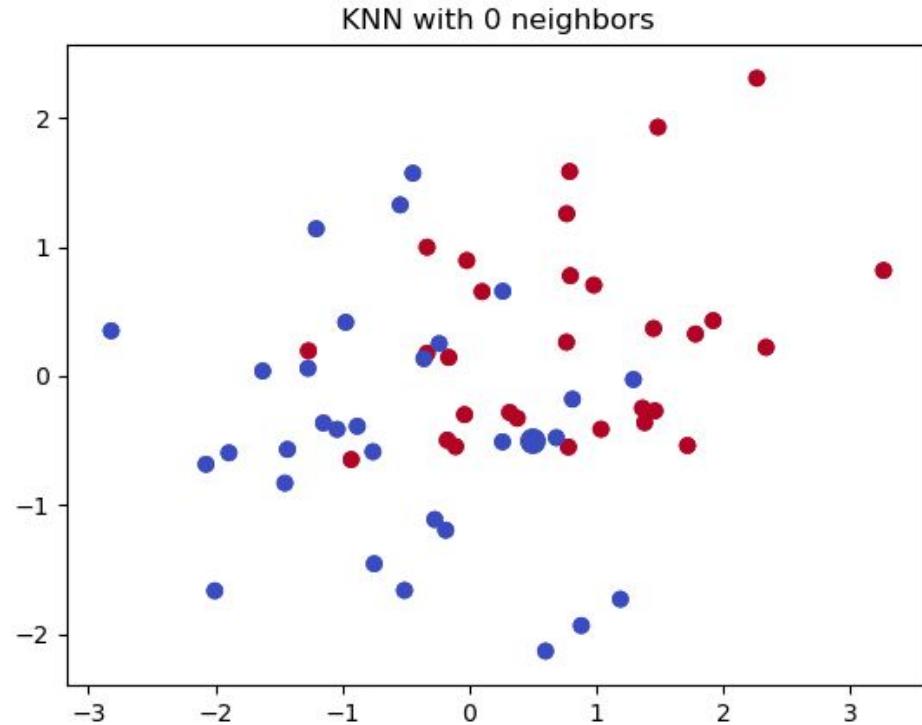
Test 1 Right

Tested Positive

Tested Negative

Modelos de vecinos

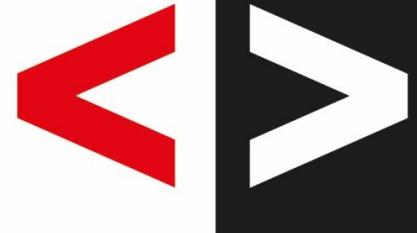
k-Nearest Neighbours: Trata de buscar los K puntos / vecinos más cercanos a un punto concreto para poder inferir su valor.



Criterios de éxito

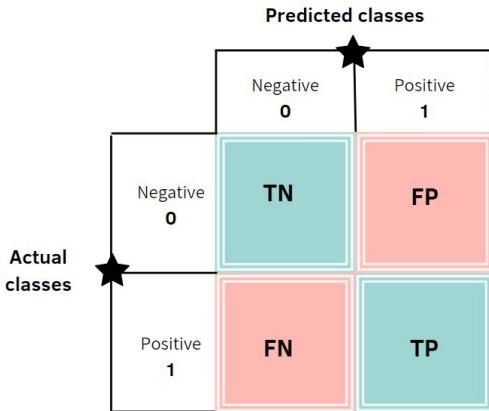
Queremos:

- **Maximizar la precisión** de la predicción
- **Minimizar el tamaño de la hipótesis**
- **Maximizar la adecuación** de la hipótesis a los datos de entrada
- **Maximizar la interpretabilidad** de la hipótesis
- **Minimizar la complejidad** temporal de la predicción



Métricas de rendimiento

- **Matriz de confusión:** resumen de los resultados de predicción (Falsos Positivos **FP**, Falsos Negativos **FN**, Verdaderos Positivos **TP**, Verdaderos Negativos **TN**) por cada combinación 1 a 1 de clases.



		Actual		Predicted	
		Positive	Negative	Positive	Negative
Actual	Positive	23	7		
		10	60		
		Positive	Negative		

		Actual				
		Elephant	Monkey	Fish	Lion	
Actual	Elephant	25	3	0	2	
		3	53	2	3	
Actual	Monkey	2	1	24	2	
		1	0	2	71	
		Elephant	Monkey	Fish	Lion	
		Predicted				

Métricas de rendimiento

Recall: Cuántos positivos consigo detectar del total de positivos existentes (cuánto abarco)

Precisión: Qué tasa de positivos correctos vs positivos incorrectos tengo (cuanto aprieto)

F1: Promedio (media armónica) de Recall y Precisión

“

El que mucho abarca poco aprieta

”



Métricas de rendimiento

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN}$$

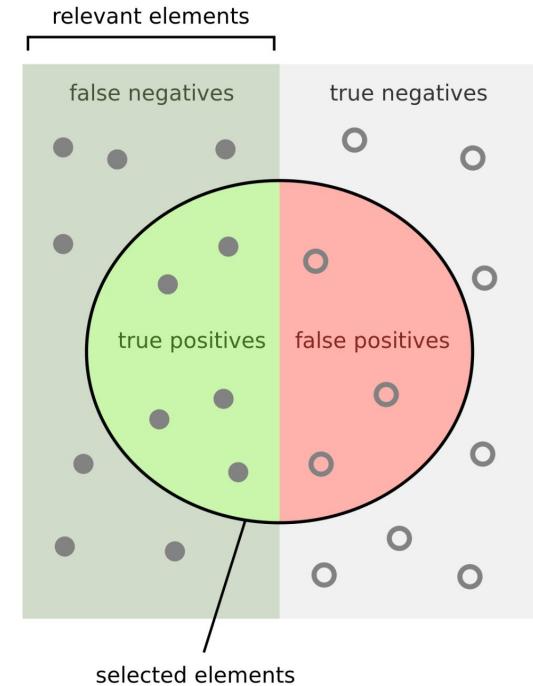
$$F1 = 2 \cdot \frac{precision \cdot recall}{precision + recall}$$

TP = True positive

TN = True negative

FP = False positive

FN = False negative



How many selected items are relevant?

$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{grey}}$$

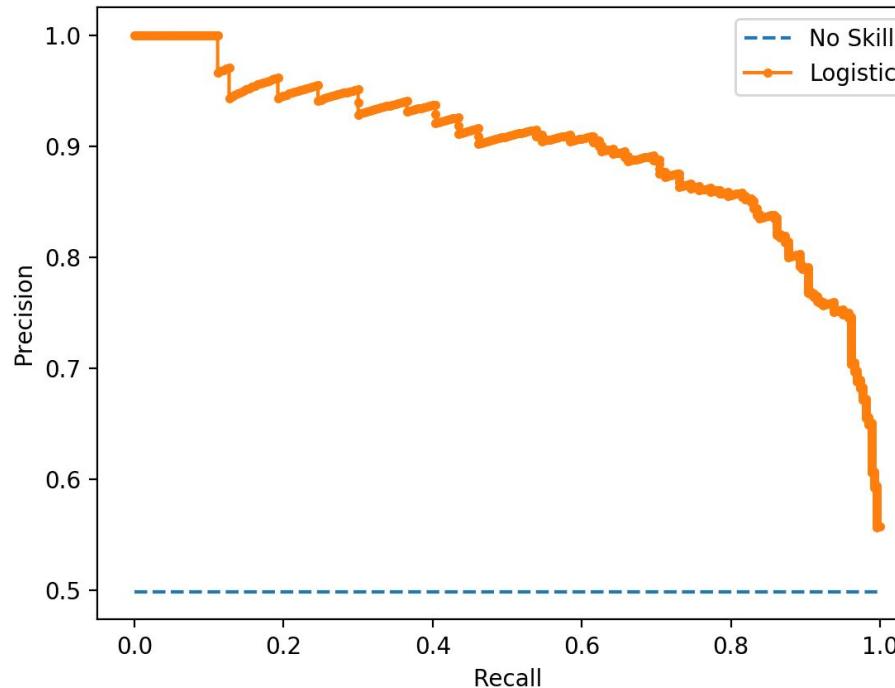
Métricas de rendimiento

Curva Precisión / Recall: Como evoluciona la precisión cuanto más abarcamos (aumentamos recall)

“

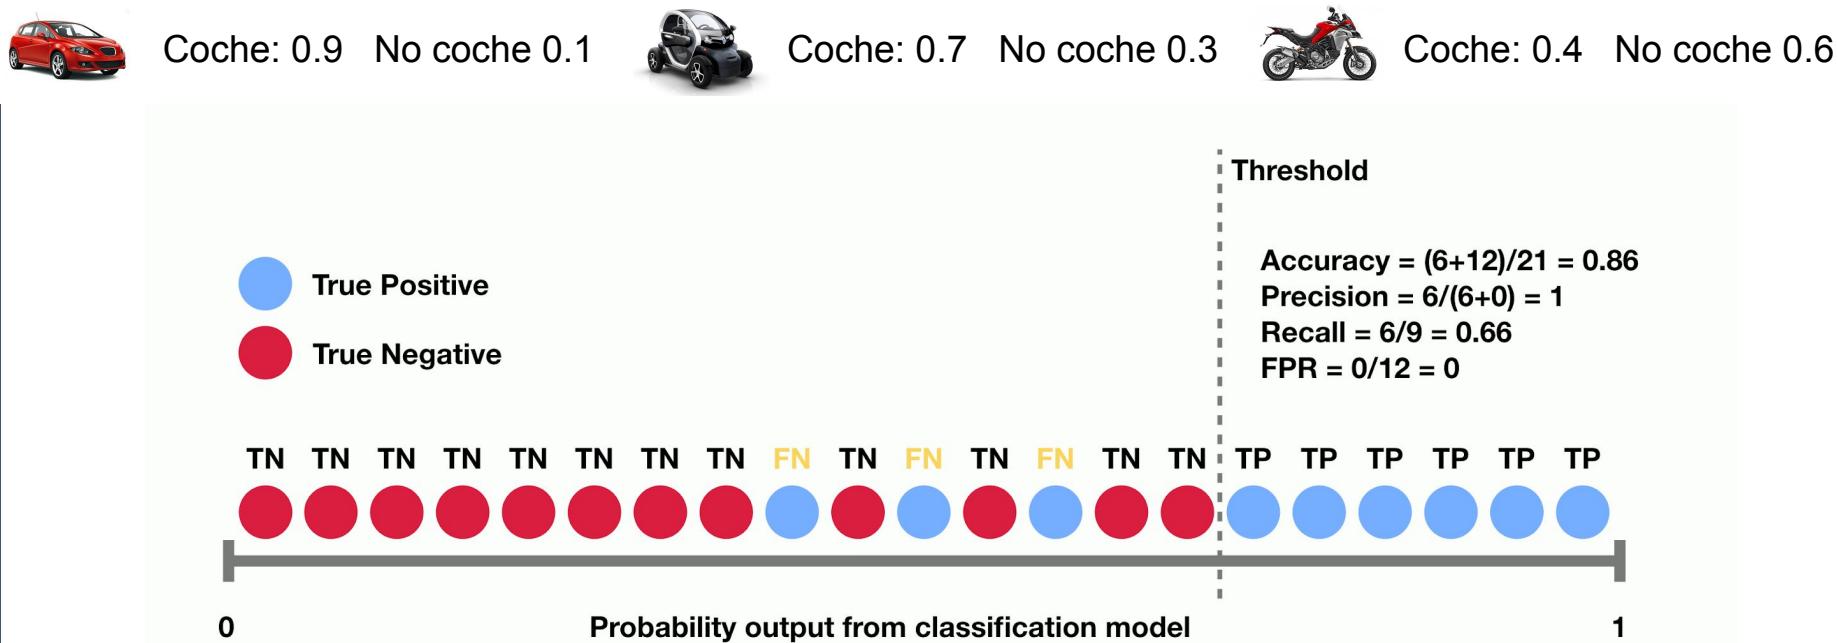
El que mucho abarca poco aprieta

”



Métricas de rendimiento

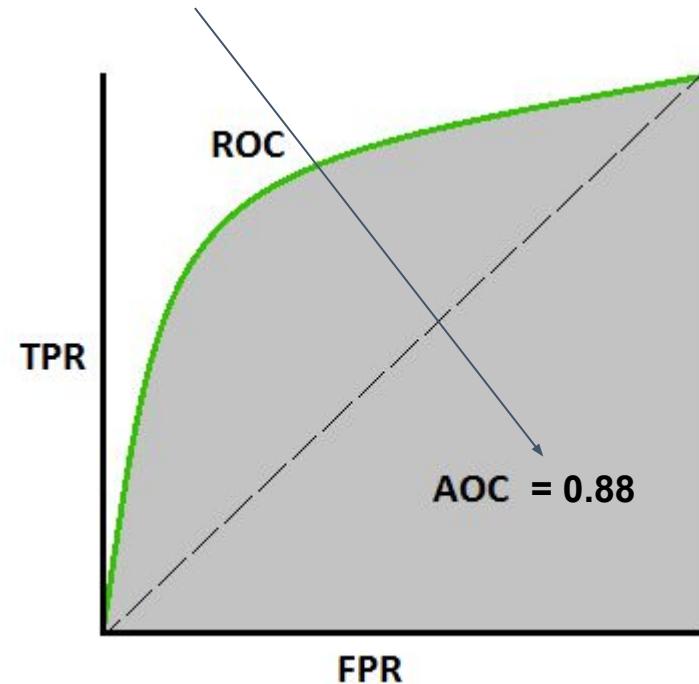
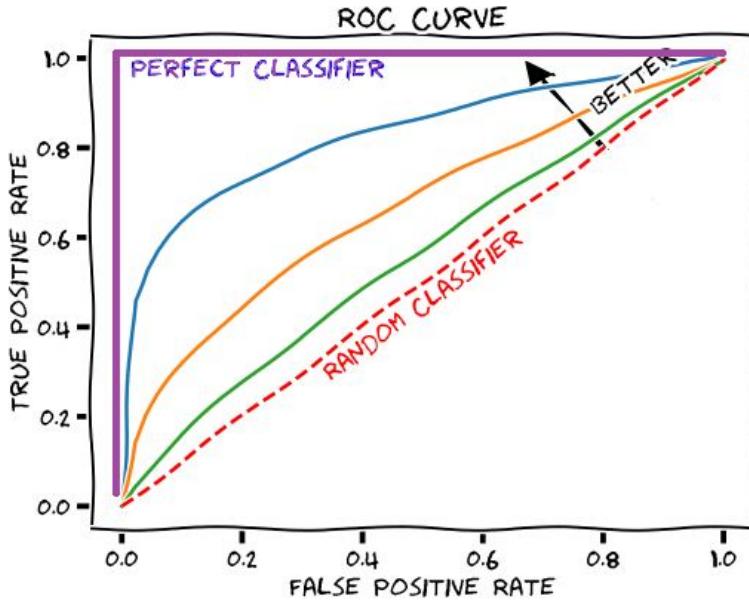
ROC: Como evoluciona la tasa Verdaderos Positivos (TP) vs Falsos Positivos (FP) en diferentes **thresholds**.



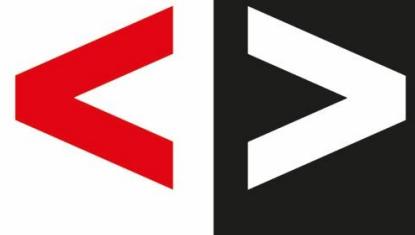
Métricas de rendimiento

ROC: Como evoluciona la tasa Verdaderos Positivos (TP) vs Falsos Positivos (FP) en diferentes **thresholds**.

AUC: Area Under the Curve. Versión numérica que cuantifica el ROC



Práctica clasificación

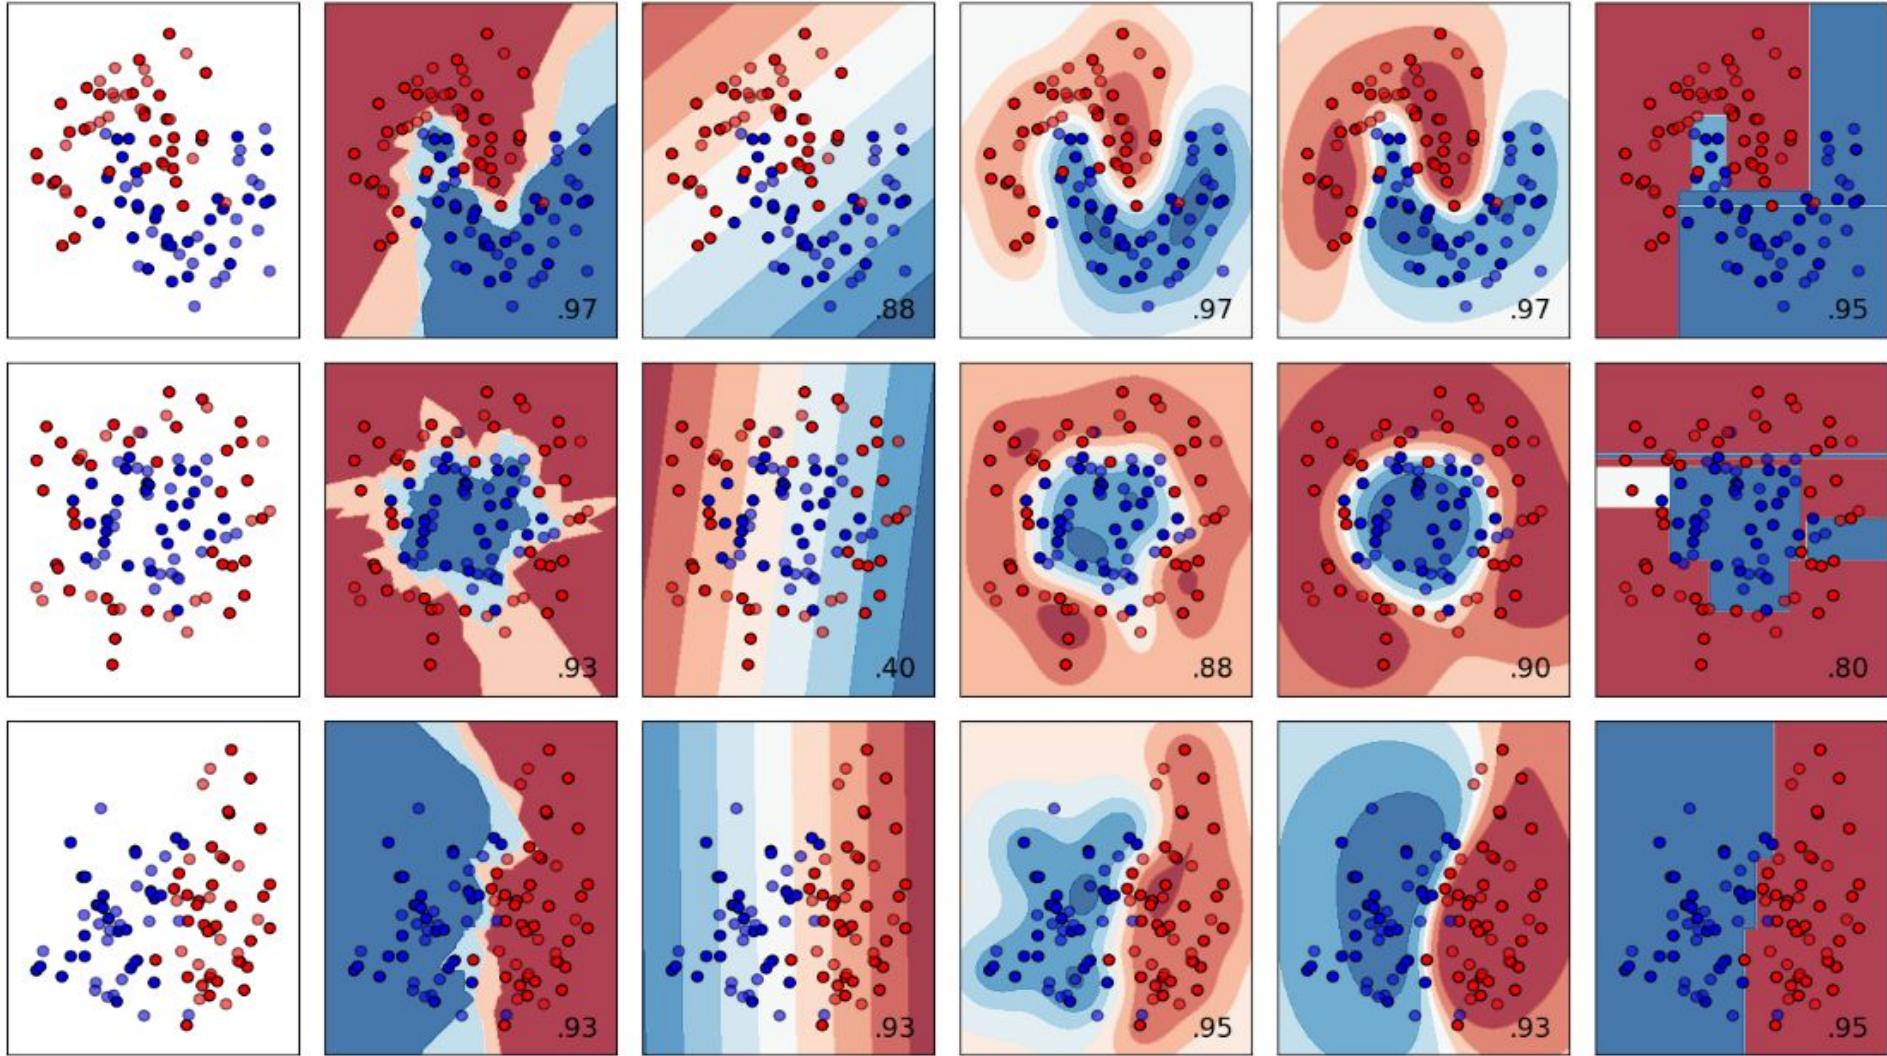


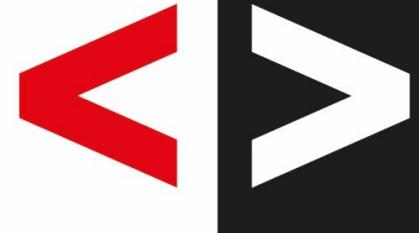
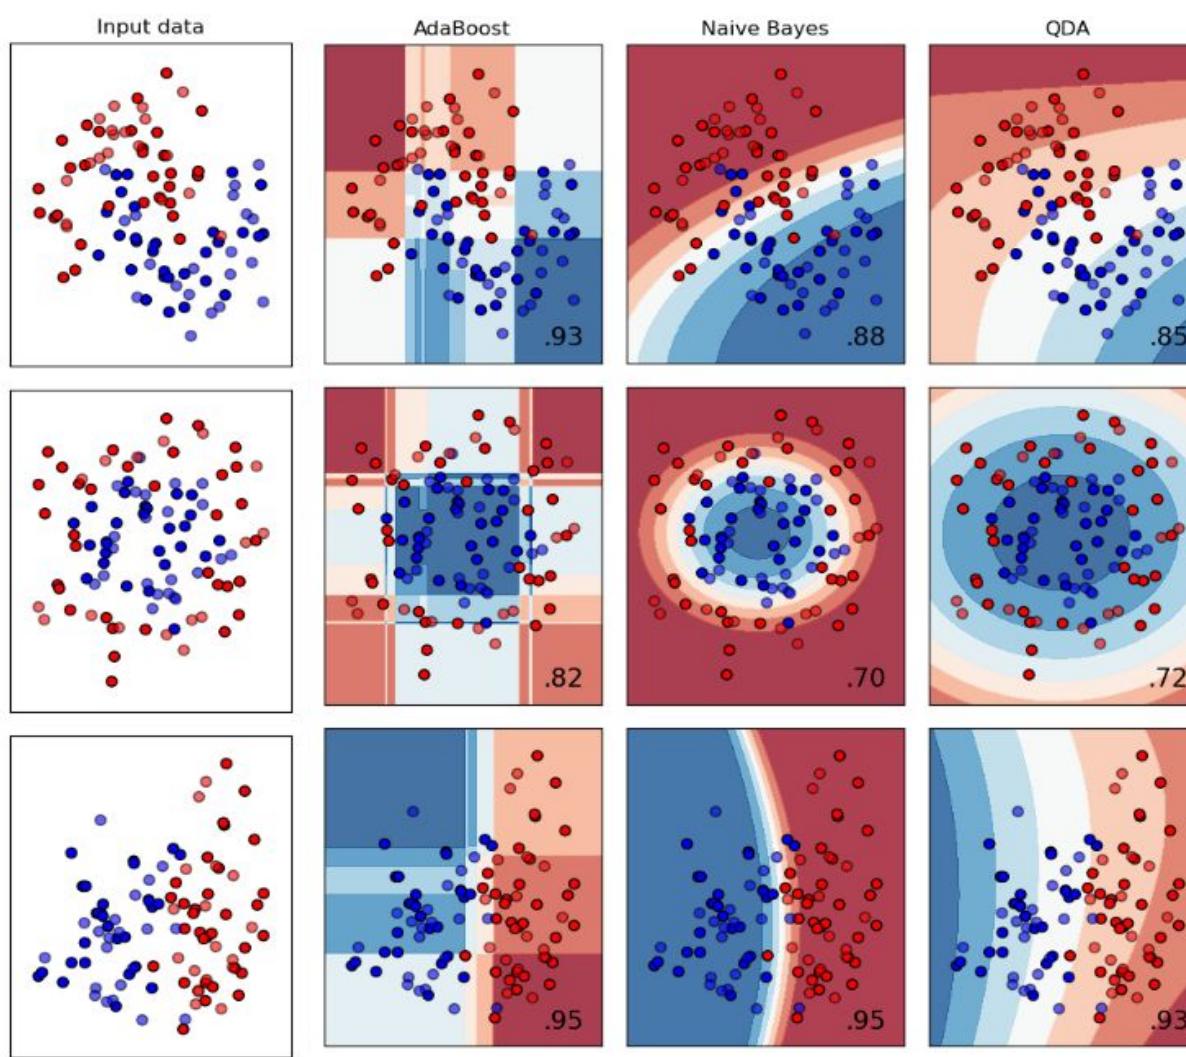
Práctica

Vamos a poner en práctica todo el conocimiento que hemos adquirido en la sesión teórica, y vamos a :

- 1) “Dibujar” nuestro propio dataset utilizando <https://drawdata.xyz/>
- 2) Completar código para crear clasificadores
- 3) Utilizar algoritmos de clasificación diferentes
- 4) Visualizar las clasificaciones
- 5) Calcular métricas
- 6) Decidir cuál es el mejor para cada caso

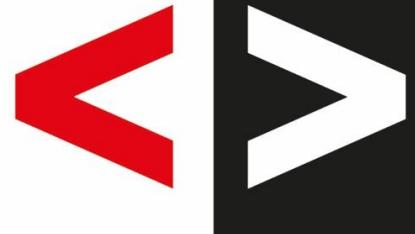


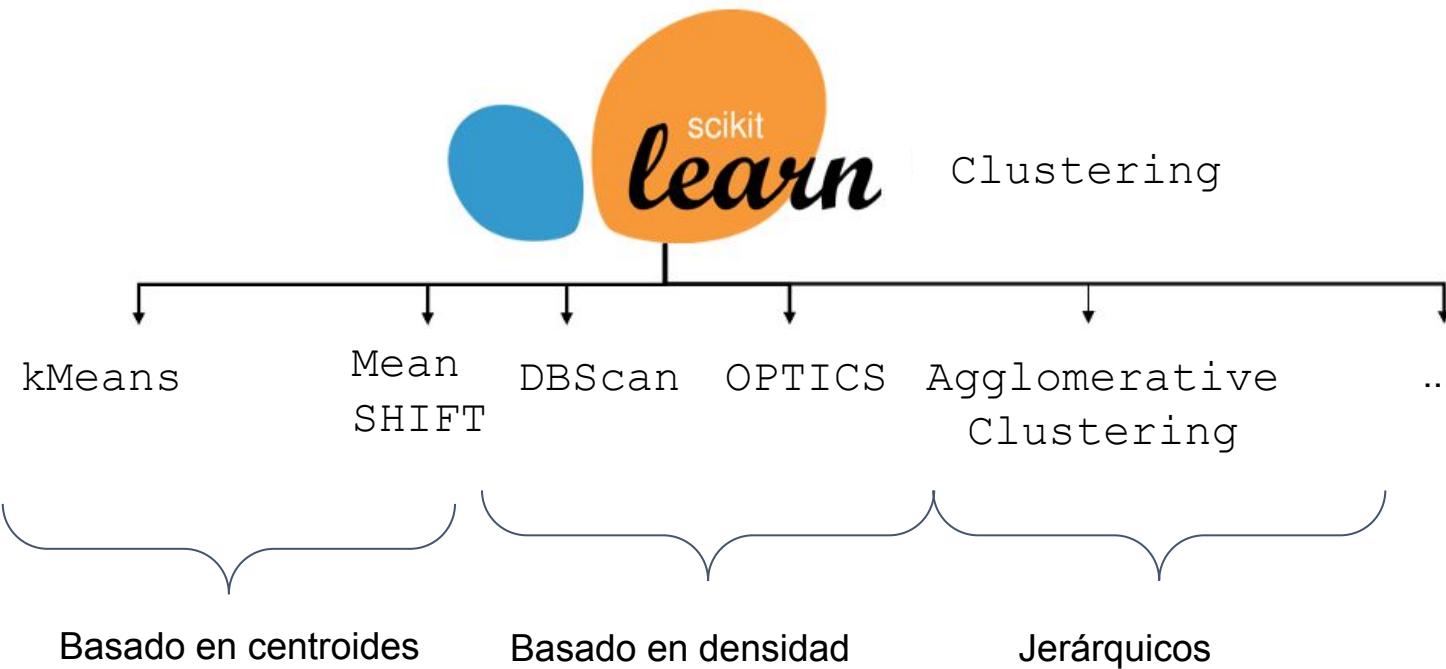




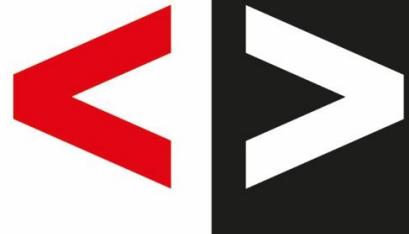
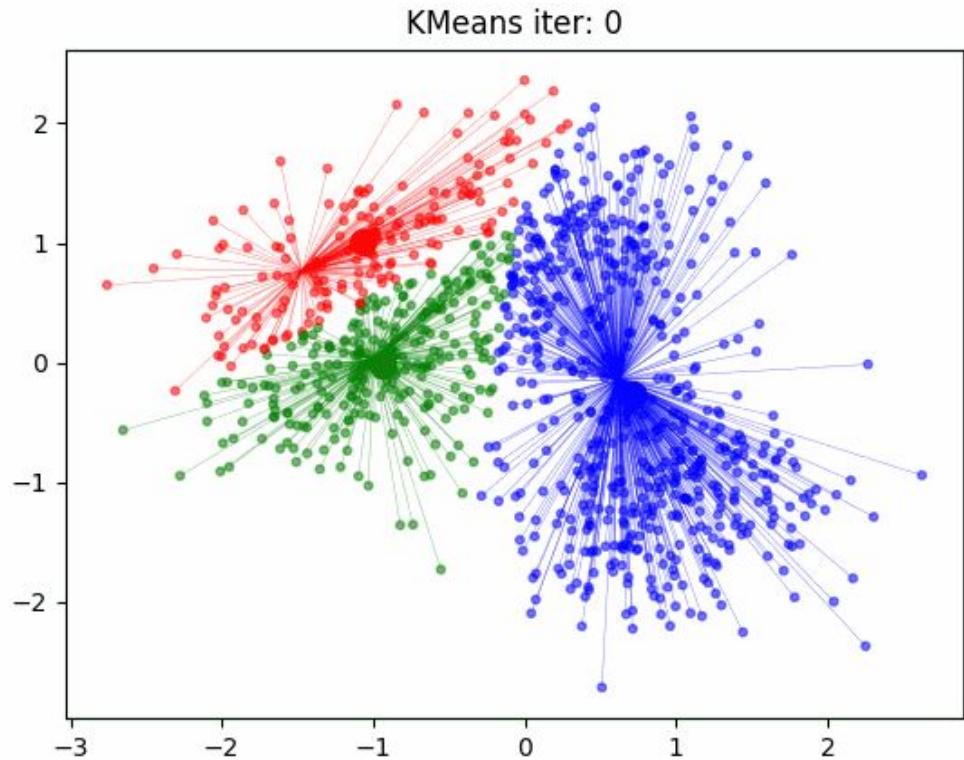
<ISDI>

Machine Learning supervisado: Clusterización

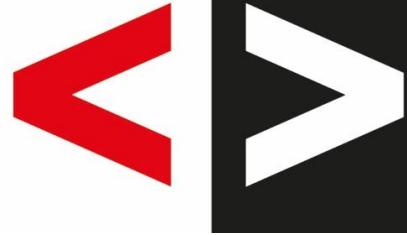
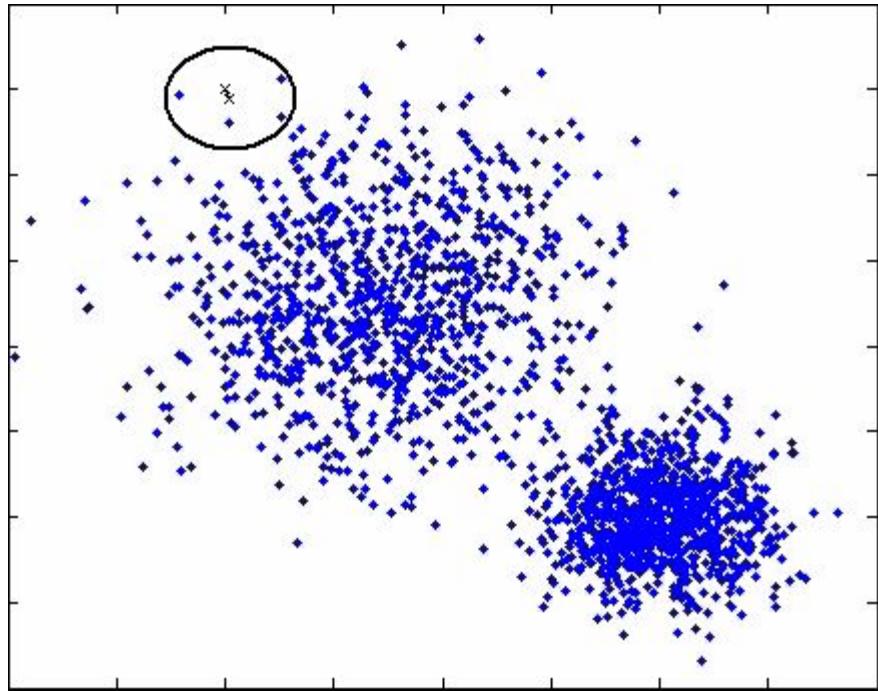




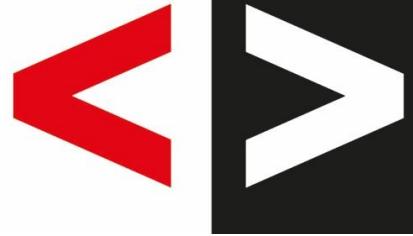
k- Means: basada en centroides



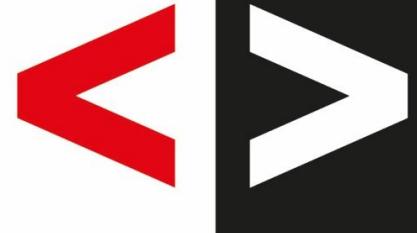
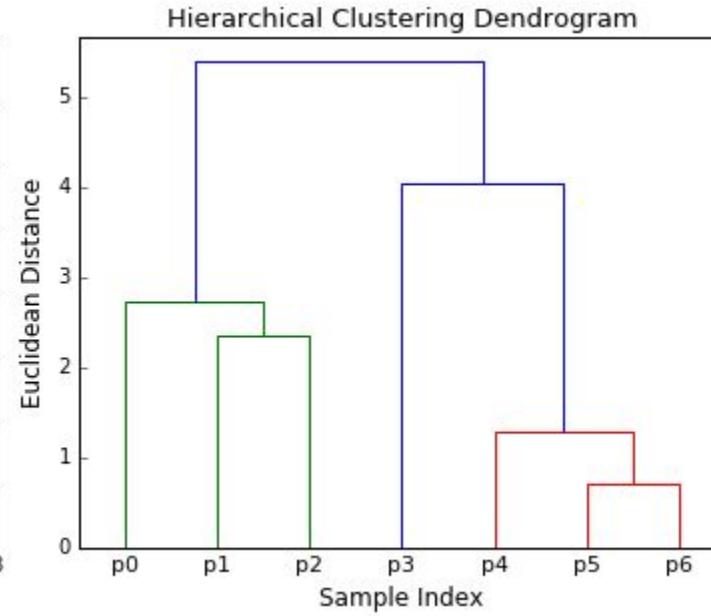
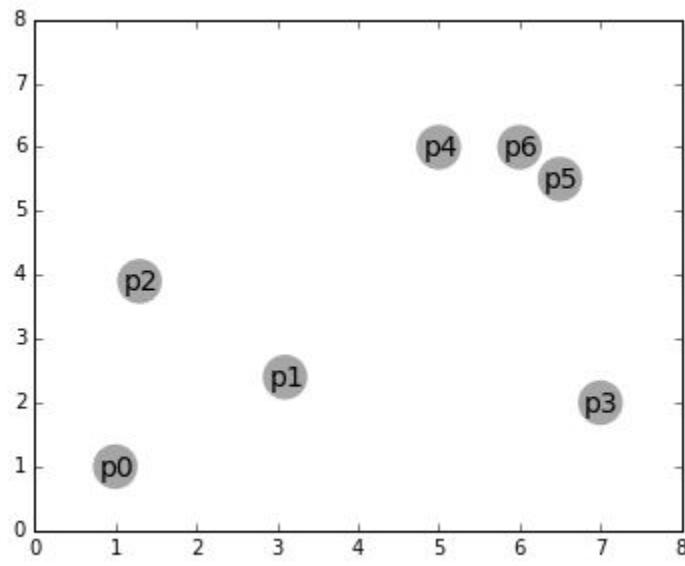
Mean Shift



DBScan: Basado en densidad



Clustering jerárquico



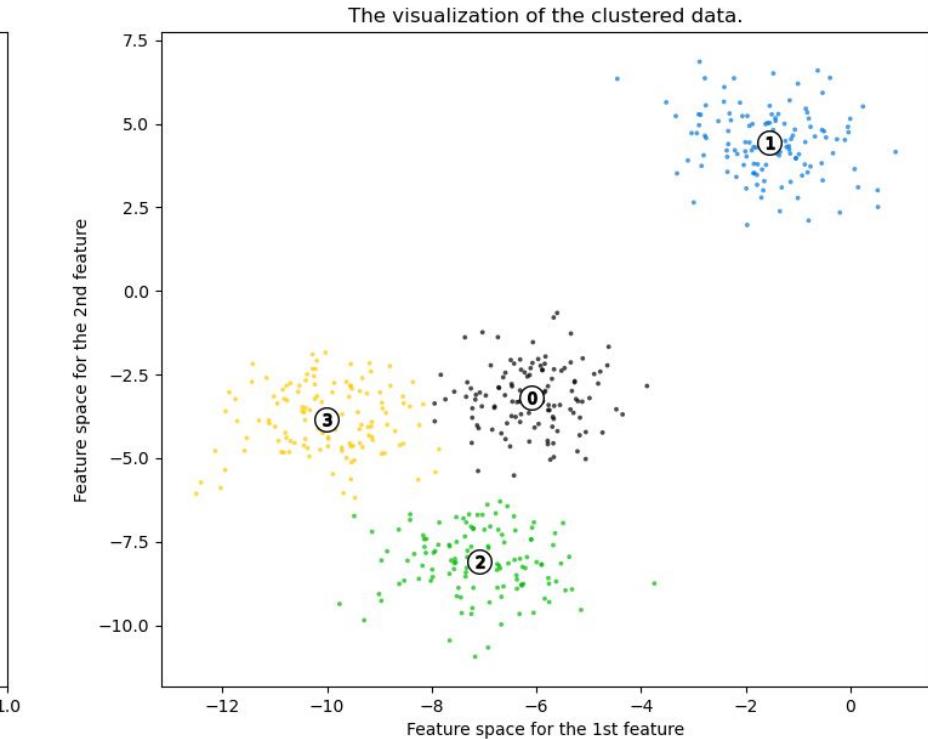
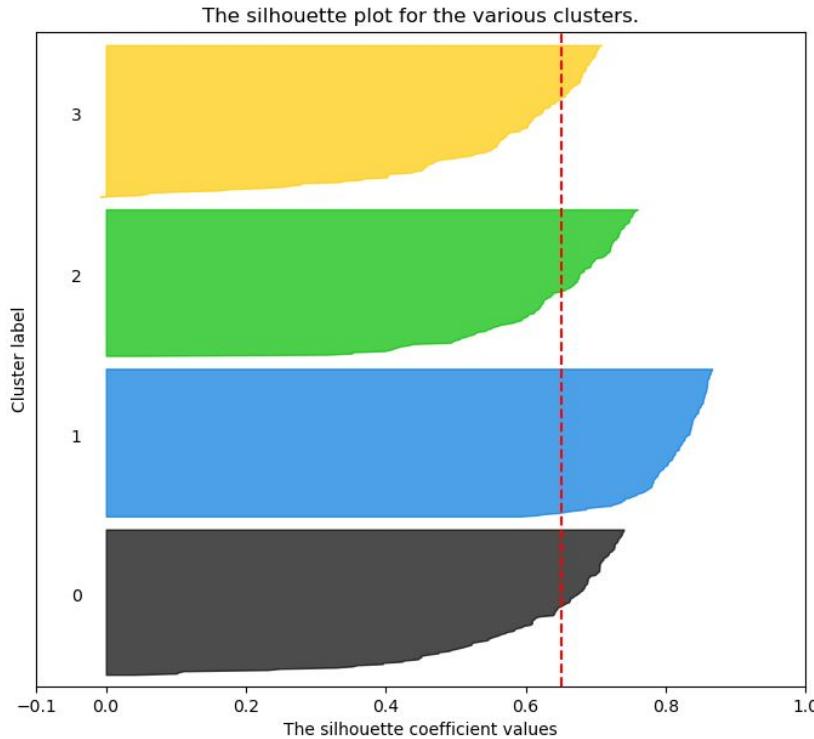
Métricas en clusterización

En regresión y clasificación (aprendizaje supervisado), **los datos están etiquetados**, y podemos calcular la diferencia con el dato esperado.

En aprendizaje no supervisado, **no hay datos etiquetados**.

Métricas: Silhouette (centroides y BIRCH)

Silhouette analysis for KMeans clustering on sample data with n_clusters = 4



Anchura: número de elementos en el cluster

Valores cercanos a 1: Alejamiento de los clústeres vecinos

Otras métricas (índices)

Sin datos etiquetados

Davis-Boulding: El score se define como la media de la similitud de cada cluster con su cluster más similar. Entendemos similitud como la relación entre las distancias dentro del grupo y las distancias entre grupos.

Los grupos que estén más separados y menos dispersos darán como resultado una mejor puntuación. Valor óptimo: 0

Con datos etiquetados

Mutual Information Score: Similitud entre la información compartida entre los clusteres y los datos etiquetados.

Valor óptimo: 1

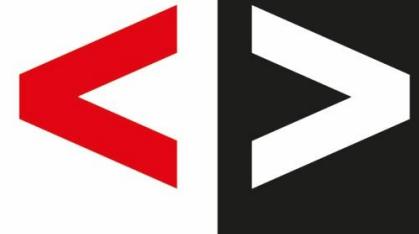
Perfect labelings are both homogeneous and complete, hence have score 1.0:

```
>>> from sklearn.metrics.cluster import adjusted_mutual_info_score
>>> adjusted_mutual_info_score([0, 0, 1, 1], [0, 0, 1, 1])
...
1.0
>>> adjusted_mutual_info_score([0, 0, 1, 1], [1, 1, 0, 0])
...
1.0
```

If classes members are completely split across different clusters, the assignment is totally in-complete, hence the AMI is null:

```
>>> adjusted_mutual_info_score([0, 0, 0, 0], [0, 1, 2, 3])
...
0.0
```

Práctica clusterización

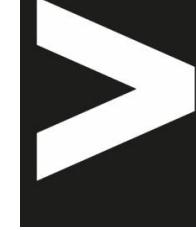


Práctica

Vamos a poner en práctica todo el conocimiento que hemos adquirido en la sesión teórica, y vamos a :

- 1) “Dibujar” nuestro propio dataset utilizando <https://drawdata.xyz/>
- 2) Completar código para crear clusterizadores
- 3) Utilizar algoritmos de clusterización diferentes
- 4) Visualizar las clasificaciones
- 5) Calcular métricas
- 6) Decidir cuál es el mejor para cada caso





< ISDI >

